

index

Symbols

[interval] [count] 55
[tcpActiveOpens] 58
[tcpAttemptFails] 58
[tcpEstabResets] 58
[tcpInErrs] 58
[tcpInSegs] 58
[tcpOutRsts] 58
[tcpOutSegs] 58
[tcpPassiveOpens] 58
[tcpRetransSegs] 58
/api/does-not-exist endpoint 259
/api/v1/ endpoint 32–33
/healthz endpoint 315
.legacy_server command 172
.org/my subfolder structure 208
/search page 236
%CPU field 61
%ifutil field 56
~/src/examples/app command 234

A

ab command 91, 95, 97, 161–164, 189, 191, 194, 236–237, 239
accept syscall 182, 184
access syscall 179
active/s field 58
addTransformer method 213, 220
affinity 321
Agent class 212, 214

agentmain method 227
ALL keyword 156
alpine image 114
Alpine Linux 114
anti-affinity 321, 332
Apache Bench 31
Apache2 85
apache2 package 361
apache2-utils package 360
app=goldpinger label 278, 291, 296
AppKiller assault 226
application layer 75–79
 BCC suite and Python 77–79
 cProfile module 76–77
application-level fault injection 228–245
application vs.
 infrastructure 243–244
experiments
 failing requests 241–243
 Redis latency 235–240
 scenario 229–235
.apply(this, arguments)
 method 254
@app.route(231
assaults 226
atexit(3) function 177
attach_chaos_if_enabled function 238–239
automated monitoring systems 79–83
 Grafana 80–83
 Prometheus 80–83
autonomous mode 304
availability 326

available column 60
available field 61

B

bandwidth toxic 294
bash program 134
bc command 68–70, 74
BCC (BPF Compiler Collection)
 BPF and 185–187
 Python and 77–79
benchmarking 5
biotop tool 53–54
biotop-bpfcc command 54
black boxes 22
blast radius 19, 36, 349
block I/O (block input/output)
 devices 50–54
 biotop tool 53–54
 df tool 51–52
 iostat tool 52–53
BPF (Berkeley Packet Filter) 53, 185–187
-bpfcc 53, 185
bpfcc-tool package 54
br-b1ac9b3f5294 interface 152
bridge mode 150
bridge option 150
brk syscall 179
buffers 60
bytecode 207–215
 -javaagent 211–215
 reading 208–210
Byteman 223–225
 installing 223
 using 223–225

Byte-Monkey 225–226
 installing 225
 using 225–226

C

CACHE_CLIENT variable 237–239, 241
 capabilities 155–157
`--cap-add` flag 156
`--cap-drop ALL` flag 156
`--cap-drop` flag 156
CAP_KILL 155
cap_sys_chroot 155
CAP_SYS_TIME 155
cat /proc/loadavg command 49
`.catch` handler 259
 catch method 259
 cereal_killer.sh script 35–36, 38
cgcreate 71, 136
cexec 71, 136
cgroup.procs 139
cgroups tool
 implementing simple containers with 146–149
 killing processes in namespace 135–140
 chaos engineering 1–15, 345, 350–358
 buy-in 349–351
 game days 350–351
 management 349–350
 team members 350
 defined
 what it is 2–3
 what it isn't 11–12
 FaaS example 13–15
 all-night investigation into problem 13–14
 four steps 15
 overview of 13
 postmortem 14–15
 four steps of 6–11, 377
 experiment 11
 hypothesis 10–11
 observability 9
 steady state 10
 mindset for 346–348
 failing early vs. failing late 347–348
 failure will happen 347
 motivations for 3–6
 estimating risk and cost 3–5
 finding emergent properties 5–6

setting SLIs, SLOs, and SLAs 3–5
 testing system as a whole 5
 problems with term 378
 teams as distributed systems 351–357
 bottlenecks 355–356
 misinformation and trust 354–355
 single points of failure 353–354
 testing processes 356–357
 tool installation 359–366
 Linux tools 360–363
 Minikube 364–366
 prerequisites for 359
 source code 364
 WordPress 363–364
 tools comparison 376
CHAOS environment variable 238, 242
 Chaos Monkey for Spring Boot 226–227
 chaos network 152
 chaos pizza 382–384
 ingredients 383
 preparation 383–384
 chaos proxy 299–300
 ChaosClient class 238
CHAOS_DELAY_SECONDS variable 238
 chaos-engineering recipes 379–384
 chaos pizza 382–384
 SRE burger 379–382
 ChaosMachine 227
 chroot tool
 implementing simple container 117–119
 uprooting processes 114–117
classfileBuffer 212
ClassFileTransformer interface 212
ClassInjector class 218–220
className 212
ClassNode class 218
ClassPrinter class 212–213
ClassReader instance 217
CLIENT_PORT_OVERRIDE environment variable 295
clock_nanosleep syscall 180–181
close syscall 189–193
 analysis 192–193
 implementation 191–192
 steady state 191
CLOSED state 58
CLOSE-WAIT state 58
 cloud 375
 cloud layer 318–323
 availability zones 319–321
 cloud provider APIs 319–321
 taking VMs down 321–323
ClusterRole 276
 clusters 324–342
 control plane 325–332
 `etcd` 326–329
 `kube-apiserver` 329–330
 `kube-controller-manager` 330–331
 `kube-scheduler` 332
 Docker, and container runtimes 335–338
 networking 338–342
 ingress networking 342
 pod-to-pod networking 339–340
 service networking 340–342
 pause container 333–335
 setting up using Minikube 272
 starting 272–274
Cmd 165
CNCF (Cloud Native Computing Foundation) 269
CNI (Container Networking Interface) 339
coll/s field 57
 command column 125
 consensus 326
 container image format 112
 Container Runtime Interface (CRI) 336
 container runtimes 112, 335, 343
ContainerCreating 286
containerd 336
CONTAINER_PID 154
 containers
 Docker and 110–112, 335–338
 implementing with
 cgroups 146–149
 implementing with chroot tool 117–119
 implementing with namespaces 133–135
 network slowness for 161–165
 experiment implementation 162–165
 Pumba 161–162

containers (*continued*)
 one container preventing another from writing to disk 119–124
 pause container 333–335
 virtual machines and 107–110
 containers, defined 104
 Content-type header 236
 control container 119, 123–124
 Control group (cname) 125
 control groups 71, 135
 control plane 271
 cProfile module 76–77
 CPU (central processing unit) 66–72
 mpstat P ALL 1 tool 69–70
 stress command 141–143
 top tool 67–69
 cpu controller 136
 --cpu flag 143
 cpu type 147
 cpu,cpuacct 136, 149
 cpu.cfs_period_us 137, 143
 cpu.cfs_period_us,
 cpu.cfs_quota_us 137
 cpu.cfs_quota_us 137, 143
 --cpu flag 141
 cpu.shares 137, 143
 CRI (Container Runtime Interface) 336
 CRI-O 336
 cumtime 76
 curl command 33, 97
 curl package 360

D

daemon restarts 166
 DaemonSet 279
 databases 84–100
 testing in production 98–100
 WordPress weak link
 example 86–98
 overview of 85–86
 slow connection
 experiment 92–98
 slow disks experiment 87–92
 deployment type 330
 DEV keyword 55
 df-h command 123
 df tool 51–52, 60
 dispatchEvent method 257

dmesg tool 27, 49–50
 DNS (Domain Name System)
 server 5
 Docker 103–168
 advanced networking 167
 blocking syscalls with seccomp 196–197
 chroot tool
 implementing simple container 117–119
 uprooting processes 114–117
 container runtimes 335–338
 containers 110–112
 daemon restarts 166
 experiments
 CPU usage 141–143
 killing processes in different PID
 namespace 129–140
 network slowness for containers with Pumba 161–165
 one container preventing another from writing to disk 119–124
 RAM overuse 143–149
 history of 106–110
 emulation, simulation, and virtualization 106–107
 virtual machines and containers 107–110
 namespaces 127–129
 isolating processes with 124–127
 networking 150–157
 capabilities 155–157
 seccomp 157
 security 167–168
 slow app problem 104–105
 solving problem 158–161
 storage for image layers 166–167
 docker command-line client 112
 Docker Hub 112
 docker images command 141
 docker inspect command 116
 docker inspect firstcontainer command 116
 docker network
 subcommand 150
 docker ps command 128
 Docker Registry 112

docker run 156
 docker stack deploy 158
 docker stack ls command 160
 docker utility 112
 docker0 bridge interface 150
 dockerd 112
 docker.io package 360
 Domain Name System (DNS)
 server 5
 down toxic 294
 downtime 4
 -driver flag 273
 dtrace tool 75

E

-e inject flag 190
 -e inject option 190–191
 e2e (end-to-end) tests 5, 241
 EACCES error 190
 eBPF (extended Berkeley Packet Filter) 53, 185
 ECHO_REQUEST datagram 94
 ECHO_RESPONSE 94
 EDEV keyword 56
 EINVAL (Invalid argument) 183
 emergent properties 5–6
 emulation 106–107
 Endpoints field 281
 end-to-end (e2e) tests 5, 241
 Entrypoint 165
 error event 257
 ERRORS section 192
 ESTABLISHED state 58
 estres/s field 58
 etcd 326–329
 ETCP keyword 57
 eth0 interface 55–56
 events 256
 example service 342
 Example1 class 209
 Example1.java program 208
 exception assault 226
 exec variants 74, 172
 execsnoop tool 74
 execve 179
 exit codes 23
 _exit(2) syscall 177
 exit(3) glibc syscall 177
 exit_group 181
 export CONTAINER_ID 138, 145
 extended Berkeley Packet Filter (eBPF) 53, 185

F

f flag 24
FaaS (FizzBuzz as a Service)
 example 13–15, 19–42
 all-night investigation into
 problem 13–14
 blast radiiuses 36–37
 four steps 15, 29–36
 experiment 15, 34–36
 hypothesis 15, 34
 observability 15, 33–34
 steady state 15, 34
 Linux forensics 22–28
 exit codes 23
 killing processes 24–25
 Out-of-Memory Killer 26–28
 overview of 13
 postmortem 14–15
 scenario 21–22
 solution 38–41
 source code 20–21
faas001_a 32
faas001_b 32, 41
 Failed requests 34
 failed state 313–314
 failing early 348–350
 failing late 348
 failing requests
 experiment 241–243
 execution 243
 implementation 242–243
 plan 241–242
fail_timeout parameter 33
 failure container 121–124
 falllocate script 120
 fault injection
 application-level 228–245
 application vs.
 infrastructure 243–244
 failing requests
 experiment 241–243
 Redis latency
 experiment 235–240
 scenario 229–235
 JavaScript 256–259
 JVM 201–227
 existing tools 222–227
 experiment 204–222
 scenario 202–204
 fault mode 225
 fault option 189
 fault tolerance 326
FBEE (FizzBuzzEnterprise-Edition) 202

Fetch API 259–260
 fetch method 259
 file utility 127
 filesystem bundle 337
 filesystems feature 111
 Firecracker 109, 337
FizzBuzzEnterpriseEdition
 example 202–204
 experiment idea 204–206
 experiment
 implementation 215–222
 experiment plan 206–207
 JVM bytecode 207–215
FizzBuzzEnterpriseEdition/lib
 subfolder 203
FizzBuzzEnterpriseEdition.jar
 file 203
 flanneld daemon 339
flask.request.cookies.get
 method 231
 free tool 60
 freegames package 363
fstat 179
fsync 190
Ftrace 188
 full virtualization 107
 full-stack Python
 development 351
 fuzzing 6

G

game days 350
generate_legacy.py script 171
get command 274
get method 237–238
get_insterests function 243
get_interests function 237,
 241–242
getpcaps command 155–156
getpid syscall 196–199
getstatic instruction 210
 ghost container 160
git package 360, 364
glibc 176
 GNU C Library 176
 Goldpinger 268
 creating YAML files 278–280
 deploying 280–284
 goldpinger-chaos 300
 goldpinger-clusterrole
 ClusterRole 277
 goldpinger-rbac.yaml file 280
 goldpinger-serviceaccount
 ServiceAccount 277

goldpinger.yml file 280
 Grafana 80–83
 GraphDriver section 116
.GraphDriver.Data.MergedDir
 path 116
 Greasemonkey 261
grep 36
 groups 136
gVisor 109, 337

H

-h argument 60
-h flag 189, 299
 hardware interrupts (hi) 67
 hardware virtualization 107
 hardware-assisted
 virtualization 109
-hdd n option 89
-hdd option 89
head -n1 285
 Hello chaos message 214, 260
-help command 274
 hi (hardware interrupts) 67
host option 150
hostname -I command 153
-human option 50
 Hyper-V Requirements 366
 hypervisors 108
 hypothesis
 FaaS example 15, 34
 forming 10–11

I

-i flag 116
 ICANT project 267
 ICMP (Internet Control Message Protocol) 94
id (idle time) 67
if statement 240
 image layer storage 166–167
 Image Specification 337
imagePullPolicy: Always 318
 Inception-style reality 110, 156
index function 230–231
info endpoint 249
 ingress networking 342
 ingress type 342
 input/output (block I/O)
 devices 50
 instrumentation package 211,
 213
 integration tests 5, 241
 internal.jdk package 220

Internet Control Message Protocol (ICMP) 94
 invokestatic instruction 216, 218
 invokestatic JVM
 instruction 216
 invokevirtual instruction 210
 IOException 205–206, 215
 iostat tool 52–53, 89
 ip addr command 151
 ip command 152
 ipc (Interprocess Communication) 125
 iseg/s field 58
 isegerr/s field 58

J

java command 209
 Java Management Extensions (JMX) 226
 -javaagent argument 214, 223, 225
 javaagent argument 219–220, 223
 -javaagent package 211–215
 -javaagent parameter 224
 javac command-line tool 209
 javacalls 204
 java.io.IOException 224
 java.io.PrintStream 210
 java.lang.instrument
 interface 201, 207
 java.lang.instrument
 package 211–212, 222, 227
 java.lang.System class 210
 javap -c org.my.Example1
 command 209
 javap tool 209
 JMX (Java Management Extensions) 226
 jQuery 250
 JS (JavaScript) 246–262
 experiments
 adding failure 256–259
 adding latency 251–256
 Fetch API 259–260
 Greasemonkey and
 Tampermonkey 261
 pgweb 247–248
 implementation
 details 249–250
 scenario 247–250
 throttling 260–261
 - json flag 126

JVM (Java Virtual Machine)
 201–227
 existing tools 222–227
 Byteman 223–225
 Byte-Monkey 225–226
 Chaos Monkey for Spring Boot 226–227
 experiment 204–222
 bytecode 207–215
 finding right exception to throw 205–206
 idea 204–206
 implementation 215–222
 injecting code 216–220
 plan 206–207
 scenario 202–204

K

Kata Containers 109, 337
 kernel space 173
 kill command 22, 24–25, 28, 36, 172, 284
 KILL signal 25
 killer_while.sh script 38, 41
 killing processes 24–25
 in different PID
 namespace 129–140
 cgroups tool 135–140
 implementing simple containers with namespaces 133–135
 Out-of-Memory Killer 26–28
 known unknowns 45
 kube-apiserver 329–330, 333, 339
 kube-apiserver component 325, 330, 343
 kube-apiserver, kube-controller-manager 330
 kube-cloud-manager component 325
 kube-controller-manager 330–331
 kube-controller-manager component 325, 330, 332, 343
 kubectl 284–285, 314, 365
 kubectl -help command 274
 kubectl apply command 314, 325, 332
 kubectl apply -f goldpinger-chaos.yml command 298
 kubectl apply -f goldpinger-rbac.yml command 280, 285
 kubectl command 271, 283, 306–308, 317
 kubectl config file 305, 307
 kubectl configuration file 289
 kubectl cp 290
 kubectl delet 314
 kubectl exec 290
 kubectl get -help 274
 kubectl get command 285
 kubectl get pods -watch command 285–286, 313
 kubectl get pods -A command 273
 kubectl get pods command 282, 285, 307
 KubeInvaders 289
 kube-proxy component 341, 343
 Kubernetes 265–302, 324–344
 anatomy of cluster 324–342
 control plane 325–332
 Docker, and container runtimes 335–338
 networking 338–342
 pause container 333–335
 automating experiments 303–323
 cloud layer 318–323
 ongoing testing and service-level objectives 311–318
 PowerfulSeal 303–311
 history of 269–270
 key components 343
 overview of 268–272
 porting onto 266–268
 Goldpinger 268
 project documentation 267
 setting up cluster 272–274
 starting cluster 272–274
 using Minikube 272
 terminology 274–275
 testing software running on 274–302
 killing pods
 experiment 284–290
 network slowness
 experiment 290–302
 running project 274–284
 Kubernetes cluster 270
 kube-scheduler component 325, 331–332, 343
 kube-thanos.sh script 285–286
 kworker 54

L

labels 278
 latency 93–95
 JavaScript 251–256
 Redis 235–240
 latency assault 226
 latency toxic 294
 latency type 309
 layers 115
 ldc instruction 210
 ldd command 118
 legacy image 197
 legacy_server binary 171, 191
 libseccomp 198–199
 libseccomp-dev package 198
 libseccomp-devel package 198
 lightweight isolation 107
 Linux 22–28, 272
 exit codes 23
 killing processes 24–25
 Out-of-Memory Killer 26–28
 Linux containers 110
 LISTEN state 58
 lo network interface 55
 load averages 49
 LowerDir 116
 lsns command 125, 127, 129,
 153
 -lwxrq pod 286

M

main method 208–209,
 213–216
 man 2 read command 176
 man 2 syscall-name 181
 man 2 syscalls command 175
 man 3 read 177
 man 8 ld.so 179
 man cgroups 140, 158
 man command 175
 man man command 174
 man namespaces 158
 man proc 49
 man ps command 132
 man sar 56
 man strace(1) 183
 man tc 93
 man top 61
 man unshare 133
 manpages package 361
 manpages-posix package 361
 manpages-posix-dev
 package 361

max_fails 33
 mean time between failure
 (MTBF) 347
 mean time to failure
 (MTTF) 319
 memory assault 226
 memory cgroup 149
 memory controller 136
 -memory flag 143
 memory utilization 332
 memory.limit_in_bytes value
 139–140
 memory.usage_in_bytes value
 139–140
 meower 160
 MergedDir 116
 Minikube
 installing 364–366
 Linux 364–365
 macOS 365–366
 Windows 366
 setting up clusters 272
 minikube service
 command 282, 299
 minikube service goldpinger
 command 282, 297, 310
 minikube start -driver 273
 minikube start command 273
 minikube stop 273
 mmap syscall 179–180

N

-n1 flag 26
 Name member 116
 namespaces
 Docker and 127–129
 isolating processes with
 124–127
 killing processes in 129–140
 cgroups tool 135–140
 implementing simple con-
 tainers with namespaces
 133–135
 nanosleep syscalls 179
 ncalls 76
 net (Network) 125
 net namespaces 153–154
 netem subcommand 162
 -network host container 150
 -network none 150
 networking
 Docker 150–157
 advanced networking 167
 capabilities 155–157
 seccomp 157
 Kubernetes 338–342
 ingress networking 342
 pod-to-pod networking
 339–340
 service networking 340–342
 network interfaces 54–59
 sar tool 55–58
 tcptop tool 58–59
 network slowness and
 PowerfulSeal 308–311
 network slowness for
 containers 161–165
 experiment implemen-
 tation 162–165
 Pumba 161–162
 slowness 290–302

NULL argument 179
null driver 150
number of nines 4

O

-o flag 132
-o name flag 285
observability 9, 15
ensuring 9
FaaS example 15, 33–34
slow app problem 43–83
application layer 75–79
automated monitoring
systems 79–83
overview of 44–45
resources 47–75
solving problem 70–72
USE method 45–47
syscalls 178–188
BPF 185–187
Ftrace 188
strace tool 178–184
SystemTap 188
OCI image 337
OOM (Out-of-Memory)
Killer 25–28
oom_dump_tasks 29
oomkill tool 65–66
--oom-kill-disable flag 146
oom_reaper 27
oom_score_adj column 27
open syscall 172, 179
openat 179
openjdk-8-jdk package 361
opensnoop tool 73–74
operand stack 210
org/my/Example1 class 214
org.agent package 217
org.agent2 package 217
org.agent2.Agent 220
org.agent2.ClassInjector
class 217
org.agent.Agent 214
org.my.Example1 class 209
orsts/s field 58
OS (operating system) 73–75
 execsnoop tool 74
 opensnoop tool 73–74
oseg/s field 58
OS-level virtualization 107
out static field 210
output method 206, 221, 224
overlay2 116
OXIPROXY_URL 299

P

-p flag 187
PACKAGE 360
passive/s field 58
PATH 294, 361, 366
pause container 333–335, 339
pending state 282, 286,
313–314
perf tool 75
pgweb 247–248, 363
 implementation details
 249–250
 installing 362
pgweb -user 247
php package 361
PID (process ID) 36, 125
--pid flag 134
pid namespace 127, 129
ping command 94
Pip dependencies 363
platform virtualization 107
podAction 307, 322
pods 274
 killing
 killing half experiment
 284–288, 306–308
 tools for 289–290
 pod-to-pod networking
 339–340
 verifying pods are ready
 313–318
pod-to-pod networking
 339–340
policy file 304
PORT environment variable
 279, 295
portability 173
postgresql package 361
PowerfulSeal 303–311
 defined 304–305
 installing 306
 killing pods experiment
 306–308
 network slowness experiment
 308–311
powerfulseal -help command
 306
powerfulseal autonomous
 --policy-file experiment1b
 .yml command 307
powerfulseal autonomous
 --policy-file experiment2b
 .yml command 310
powerfulseal command 306
powerfulseal pip package 306
powerfulseal/powerfulseal
 image 306
pquota option 124
premain class 220
premain method 212–214,
219–220
Premain-Class attribute 212–214,
219
preparing dependencies 314
println call 208
.println method 210
prio qdisc 96
--privileged flag 167
probeHTTP 314, 316
process ID (PID) 36, 125
production, testing in 98–99
profile module 76
profile.json 196
Prometheus 80–83
Promise object 259
prom.yml configuration file 80
PROT_NONE flag 180
protocol 112
PROT_READ flag 180
ps command 34, 132, 148
ptrace default leverage 337
ptrace syscall 184
Pumba 161–162, 361
pumba help 161
pumba netem command 162
Python
 BCC suite and 77–79
 installing 362
python3-pip package 360
pythonflow 75, 77–78
pythonstat 75, 77–78

Q

QA (quality assurance)
 environment 37
qdisc 93
queueing discipline 93

R

raise_rediserror_every_other_
time_if_enabled function
 242
RAM (random access memory)
 59–66
 free tool 60
 oomkill tool 65–66
 stress command 143–149

RAM (random access memory)
 (continued)
 top tool 60–63
 vmstat tool 63–65
 read syscall 172, 175, 179, 183,
 186
 read-eval-print loop (REPL) 76
 readlink command 127
 recommend_other_products
 function 231
 Redis latency experiment
 235–240
 discussion 240
 execution 239–240
 implementation 237–239
 plan 235–236
 steady state 236–237
 RedisError 241
 redis.exceptions.RedisError 242
 redis-server command 234
REFRESH_PERIOD
 variable 284
 REPL (read-eval-print loop) 76
 ReplicaSet 330
 reset function 231
 resources 45, 47–75, 274
 block I/O devices 50–54
 biotop tool 53–54
 df tool 51–52
 iostat tool 52–53
 CPU 66–72
 mpstat P ALL 1 tool 69–70
 top tool 67–69
 network interfaces 54–59
 sar tool 55–58
 tcptop tool 58–59
 OS 73–75
 execsnoop tool 74
 opensnoop tool 73–74
 other tools 75
 RAM 59–66
 free tool 60
 oomkill tool 65–66
 top tool 60–63
 vmstat tool 63–65
 system overview 48–50
 dmesg tool 49–50
 uptime tool 48–49
 response.set_cookie 231
 retruns/s field 58
 retval=<return code> argument
 190
 rkB/s field 52
 rkt 336
 --rm flag 114

--rm option 124
 RST flag 58
 runc 135
 Running 281, 286
 running state 313–314
 Runtime Specification 337
 runtimes 376
 rxcmp/s field 56
 rxdrop/s field 57
 rxerr/s field 57
 rxfifo/s field 57
 rxfram/s field 57
 rxkB/s field 56
 rxmcst/s field 56
 rxpck/s field 56

S

-S count flag 181
 -s flag 64
 sar tool 55–58
 saturation 45
 scenarios 304
SCMP_ACT_ERRNO default
 action 196
 scraping metrics 80
 search function 231
 seccomp tool
 blocking syscalls 195–199
 with Docker 196–197
 with libseccomp 198–199
 Docker networking 157
 seccomp_init function 198
 seccomp_load function 198
 seccomp_release function 198
 seccomp_rule_add function 198
 Secure Shell (SSH) 269
 security feature 111, 173
 --security-opt seccomp 157
 selectors 278
 send method 254–256
 service networking 340–342
ServiceAccounts 276
 service-level agreements
 (SLAs) 3–5
 service-level indicators (SLIs) 3–
 5
 service-level objectives
 (SLOs) 3–5
 services 275
 sessionID cookie 231
 set method 237–238
 set_session_id function 230
 setTimeout function 254
 si (software interrupts) 67

SIGFPE signal 23
 signal= sig argument 190
SIGTERM 25
 simulation 106–107
 site reliability engineering
 (SRE) 3, 44
 SLAs (service-level
 agreements) 3–5
 sleep 1 command 178
 sleep 3600 24
 sleep command 24, 178
 sleep process 138
 slicer toxic 294
 SLIs (service-level indicators)
 3–5
 SLOs (service-level
 objectives) 3–5
 slow app problem 43–83
 application layer 75–79
 BCC suite and Python
 77–79
 cProfile module 76–77
 automated monitoring
 systems 79–83
 Grafana 80–83
 Prometheus 80–83
Docker 104–105
 architecture 105
 CPU usage 141–143
 killing processes in differ-
 ent PID namespace
 129–140
 network slowness for con-
 tainers with Pumba
 161–165
 one container preventing
 another from writing to
 disk 119–124
 RAM overuse 143–149
 overview of 44–45
 resources 47–75
 block I/O devices 50–54
 CPU 66–72
 network interfaces 54–59
 OS 73–75
 RAM 59–66
 system overview 48–50
 solving problem 70–72,
 158–161
 USE method 45–47
 slow close toxic 294
 slow connection experiment
 92–98
 implementation 95–98
 latency 93–95

slow disks experiment 87–92
 discussion 91–92
 implementation 88–91
 software interrupts (si) 67
 solid-state drives (SSDs) 90
 sort -random-sort 285
 source code 20–21, 364
 SPA (single-page application) 250
 SRE ('ShRoomEee)
 burger 379–382
 assembling finished product 382
 hidden dependencies 381
 ingredients 380
 making patty 381–382
 SRE (site reliability engineering) 3, 44
 SSDs (solid-state drives) 90
 SSH (Secure Shell) 269
 st (steal time) 67
 stage=dev label 278
 stapbpf 188
 StartLimitIntervalSec 40
 Staycation 353
 steady state 9, 15
 breaking close syscall 191
 breaking write syscall 194
 defining 10
 FaaS example 15, 34
 JavaScript latency 252
 stopHost action 322
 -storage-opt size 124
 store_interests function 231, 241
 strace command 178, 181–182,
 184–186, 190–191, 194–195
 strace -h flag 190
 strace output 178
 strace tool 75, 182–183
 blocking syscalls 188–195
 breaking close syscall
 189–193
 breaking write syscall
 193–195
 overhead 183–184
 sleep command and 178–181
 stress command 70, 74, 89–91,
 141–144
 stress package 360
 string type 278
 strong isolation 107
 succeeded state 313
 sudo apt-get install bpfcc-tools
 linux-headers-\$(uname -r)
 command 185

sudo apt-get install git 360
 sudo apt-get install PACKAGE 360
 sudo command 36, 119, 128,
 178
 sudo lsns -t pid command 135
 sudo password 316
 sudo pip3 install redis 234
 sudo service postgresql start
 command 247
 sudo strace command 181
 sudo strace -C command 182
 sudo syscount-bpfcc command
 185
 swap 60–61
 sy (system time) 67
 SYN-RCVD state 58
 SYN-SENT stat 58
 SYN-SENT state 58
 syscalls 169, 172–200
 blocking with seccomp
 195–199
 with Docker 196–197
 with libseccomp 198–199
 blocking with strace
 188–195
 breaking close syscall
 189–193
 breaking write syscall
 193–195
 finding out about 174–176
 observability 178–188
 BPF 185–187
 Ftrace 188
 strace tool 178–184
 SystemTap 188
 overview of 172–178
 scenario 170–172
 using standard C library and
 glibc 176–178
 syscount command 186–187
 syscount tool 185
 syscount-bpfcc 187
 sysstat package 55, 360
 system calls. *See* syscalls
 system overview 48–50
 systemctl daemon 40
 systemctl restart 36
 systemd service 36, 38, 40, 269
 systemd unit file 40
 SystemOutFizzBuzzOutput-
 Strategy class 206, 221,
 224
 SystemTap 188

T

Tampermonkey 261
 -target flag 154
 -task flag 126
 tc (Traffic Control) 235
 tc command 94, 98, 161–162,
 165, 167, 290
 tc tool 93
 -tc-image flag 162, 165
 TCP keyword 57
 tcptop tool 58–59
 telnet 97
 Terminating state 286
 testing
 chaos engineering not
 replacement for other
 methods of 12
 in production 98–99
 of system as whole 5
 .then handler 259
 then method 259
 throttled_time 143
 throttling 260–261
 throwIOException method 215,
 217–218
 throws keyword 205
 Time (time) namespace 125
 time command 97
 time curl localhost/blog/
 command 97
 timeout toxic 294
 top tool 60–64, 67–69
 tottime 76
 toxics 294
 Toxiproxy 292–294
 toxiproxy type 309
 toxiproxy-cli command 295,
 299
 toxiproxy-cli list command 299
 toxiproxy-cli toxic add
 command 300
 transform method 211–212,
 217
 TripleAgent 227
 tv_sec argument 180
 txcarr/s field 57
 txcmp/s field 56
 txdrop/s field 57
 txerr/s field 57
 txfifo/s field 57
 txkB/s field 56
 txpck/s field 56

U

unified 136
UniK 109
union filesystem 115
union mount 115
unistd.h 176–177
unit tests 5
unknown state 313
unknown unknowns 45
unshare command 133–134, 147
UpperDir 116
uptime tool 48–49
--url flag 299
us (user time) 67
USDT (User Statically Defined Tracing) 77
USE (utilization, saturation, and errors) 45–47
User ID (user) 125
user space 173
utilization 45, 51
UTS (uts) 125

V

-verbose flag 209
VERSION command 360

vim package 360
virtual machines

 containers and 107–110
 taking down 321–323

virtualenv 306
virtualization 106–107

VM (virtual machine) image 20

vmstat command 63

vmstat -D 65

vmstat -d 65

vmstat -f 65

vmstat tool 63–65

W

wa (I/O wait time) 67

--watch flag 285, 307

watch notification

 mechanism 330

wget command 59

window global scope 254

window variable 254

window.XMLHttpRequest.proto

 type.send 254, 257

--with-dtrace 78

wkB/s field 52

WORA (write once, run anywhere) principle 207

WordPress 85

wordpress package 361

WordPress weak link

 example 86–98

configuring WordPress 363–364

overview of 85–86

slow connection

 experiment 92–98
 implementation 95–98
 latency 93–95

slow disks experiment 87–92
 discussion 91–92

 implementation 88–91

write syscall 193–195

 implementation 194–195

 steady state 194

X

-x flag 52

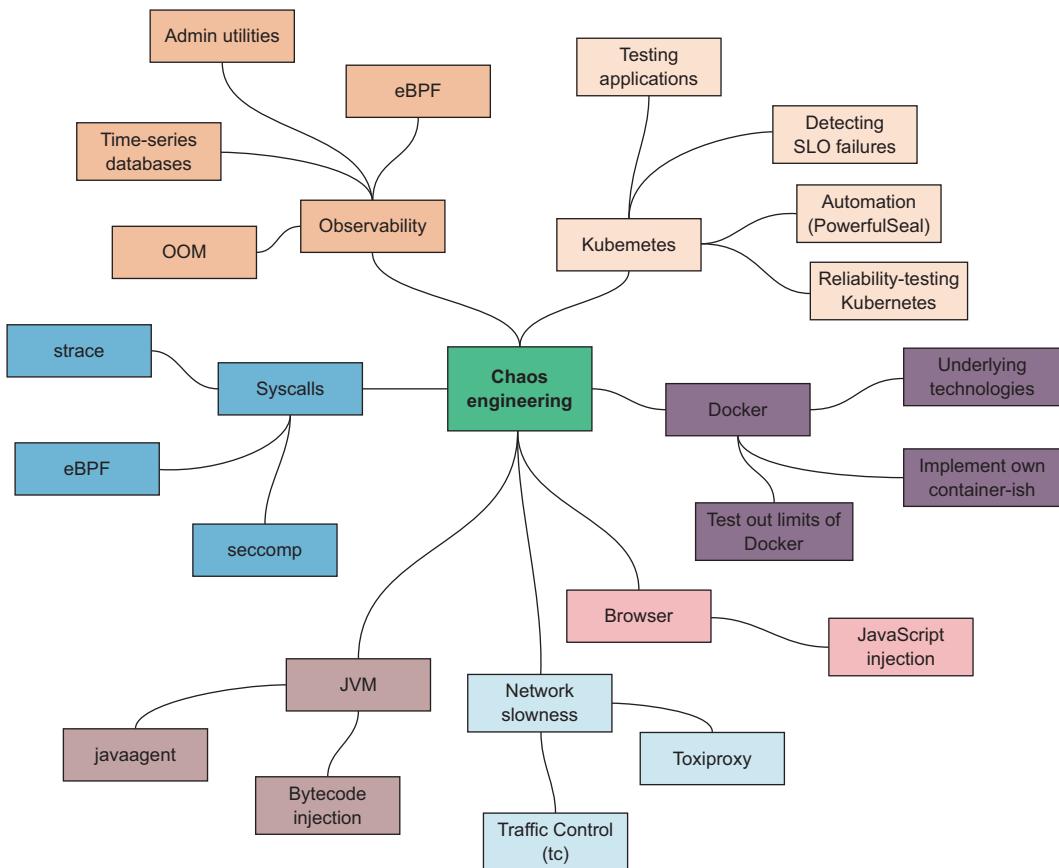
-XDignore.symbol.file flag 220

xfs filesystem 124

XMLHttpRequest 253–254, 256–257, 259–260

Y

YAML files 278–280



Chaos Engineering

Mikolaj Pawlikowski

Can your network survive a devastating failure? Could an accident bring your day-to-day operations to a halt?

Chaos engineering simulates infrastructure outages, component crashes, and other calamities to show how systems and staff respond. Testing systems in distress is the best way to ensure their future resilience, which is especially important for complex, large-scale applications with little room for downtime.

Chaos Engineering teaches you to design and execute controlled experiments that uncover hidden problems. Learn to inject system-shaking failures that disrupt system calls, networking, APIs, and Kubernetes-based microservices infrastructures. To help you practice, the book includes a downloadable Linux VM image with a suite of preconfigured tools so you can experiment quickly—without risk.

What's Inside

- Inject failure into processes, applications, and virtual machines
- Test software running on Kubernetes
- Work with both open source and legacy software
- Simulate database connection latency
- Test and improve your team's failure response

Assumes Linux servers. Basic scripting skills required.

Mikolaj Pawlikowski is a recognized authority on chaos engineering. He is the creator of the Kubernetes Chaos Engineering tool PowerfulSeal, and the networking visibility tool Goldpinger.

Register this print book to get free access to all ebook formats.
Visit <https://www.manning.com/freebook>



MANNING

\$49.99 / Can \$65.99 [INCLUDING eBOOK]

“The topics covered in this book are easy to follow and detailed. It provides a number of hands-on exercises to help the reader master chaos engineering.”

—Kelum Prabath Senanayake
Echoworx

“The book we needed to improve our system's reliability and resilience.”

—Hugo Cruz
People Driven Technology

“An important topic if you want to find hidden problems in your large system. This book gives a really good foundation.”

—Yuri Kushch, Amazon

“One of the best books about in-depth infrastructure, troubleshooting complex systems, and chaos engineering that I've ever read.”

—Lev Andelman
Terasky Cloud & Devops



ISBN: 978-1-61729-775-5



54999

9 781617 297755